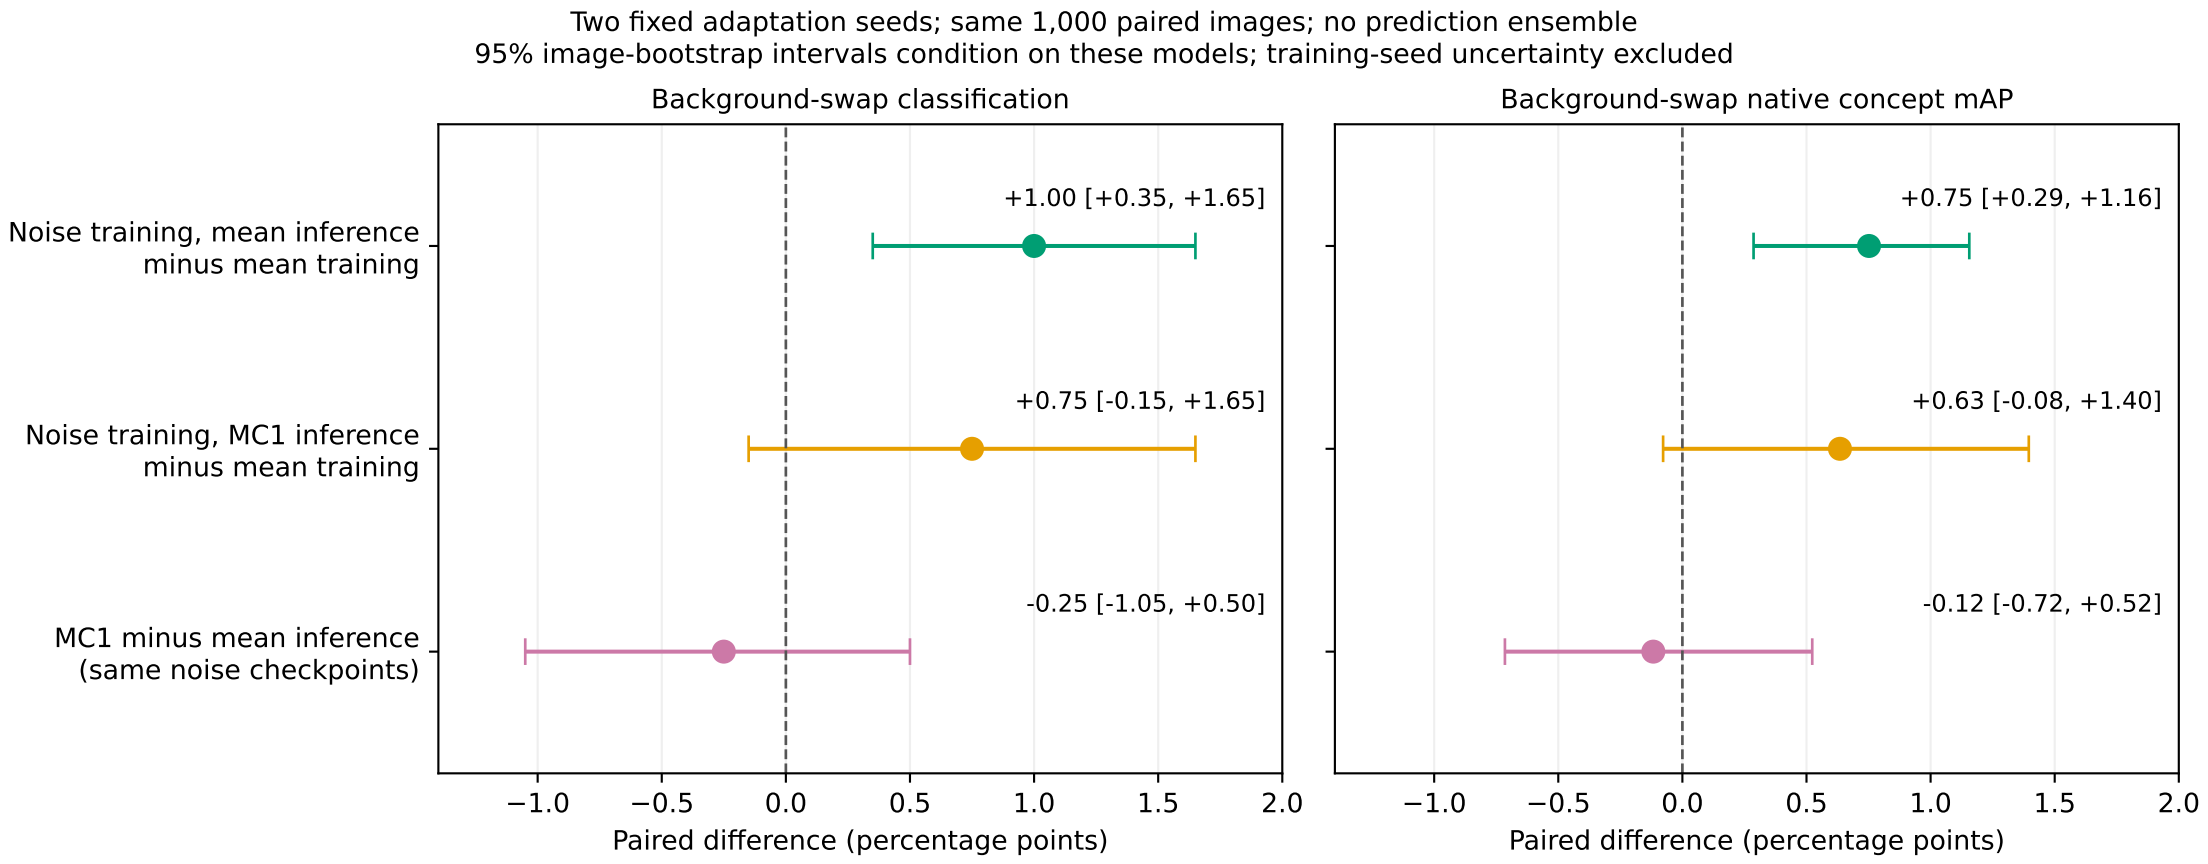
Two fixed adaptation seeds; same 1,000 paired images; no prediction ensemble
95% image-bootstrap intervals condition on these models; training-seed uncertainty excluded
Background-swap classification
Background-swap native concept mAP
Noise training, mean inference minus mean training
Noise training, MC1 inference minus mean training
MC1 minus mean inference (same noise checkpoints)
+1.00 [+0.35, +1.65]
+0.75 [-0.15, +1.65]
-0.25 [-1.05, +0.50]
+0.75 [+0.29, +1.16]
+0.63 [-0.08, +1.40]
-0.12 [-0.72, +0.52]
−1.0
−0.5
0.0
0.5
1.0
1.5
2.0
Paired difference (percentage points)